

SWE40006

Software Deployment and

Evolution

Lecture 3

Deployment -Update, Retire

Deployment - Story so far..

- Releasing a software involves:
 - Packaging and Advertising
- Contents of the package may change depending on target environment - Dev./ Test / Staging / Live
- Different packaging approaches and containers exist
- Advertising - Licensing, Support and Roadmap
 - Packaging and Advertising choices impact on product design -- developer should aim to capture these in the non-functional requirements early
- Some licenses are viral
- Activation must be fail-safe

Related Learning objectives

1. Describe the activities in software deployment and apply these in a problem context.
2. Apply knowledge of software environments to plan development and deployment.

Lecture Overview

- In this lecture we will cover,
 - **Un-installing software**
 - **Activating and De-activating**
 - **Updating and Adapting**
 - Deployment Practices and Tools
 - Multi-tier and Web Applications

Software Deployment - Deactivate




Figure 1: Activities of the Software Deployment process.

<https://www.ics.uci.edu/~andre/papers/T3.pdf>

Deactivate

- Deactivating is shutting down a software system
 - Deactivating a software system should result in returning all locks (on devices) and resources properly
- Issues:
 - Software crashes during use -- should this trigger a clean up?
 - What should be cleaned!
 - What if one part of the application needs to be stopped (i.e. database server is disconnected)
- Partial de-activation support may be needed for many systems

Software Deployment – De-Install




Figure 1: Activities of the Software Deployment Process.

<https://www.ics.uci.edu/~andre/papers/T3.pdf>

De-Install

- An ideal removal process will:
 - Remove all software components, including potentially custom drivers from the file system
 - Un-register any libraries (esp. COM components in Windows)
 - Create an archive of data files (potentially for use later or to backup) -- rarely done well
- Multi software deployments will often require deactivation of each part followed by a removal of software components

De-Install

- Removing software should be considered as an ACID* transaction
 - Should work even if power fails during the process (i.e. should recover from last state)
 - Should guarantee that no remnants are left behind (Registry keys, empty directories, hidden data directories/files etc.)
 - Currently the ACID concept is not well supported by most package managers/installers
- The entire process must be logged

*Atomicity, Consistency, Isolation, Durability

https://www.ibm.com/support/knowledgecenter/en/SSGMCP_5.3.0/com.ibm.cics.ts.productoverview.doc/concepts/acid.html

Software Deployment – Update




Figure 1: Activities of the Software Deployment Process.

<https://www.ics.uci.edu/~andre/papers/T3.pdf>

Update

- Update activity involves 'transferring and configuring' a set of components
 - Entire software can be updated as part of this as well
- Typically this process involves replacing a few components with new ones
- Update is generally performed by a separate application that is part of the installer
- Configuring a new update is similar to configuring after an install (same issues need to be considered)

Update - Transfer

- Transfer issues:
 - Is the update possible? That is, there are no conflicts with any existing components
- Configuration issues:
 - Data migration may be needed as part of the configuration step
- Other aspects:
 - Some issues and problems may not be visible until activation (e.g. If a database server was updated)

Error Handling

- You have been taught this stuff (it is all good):
 - You can and should use error codes, and also exception handling mechanisms
 - This works in most cases
- But.... here is where the real pain comes from:
 - The exception handling (or error code based recovery) mechanisms assume that you did not change the state of the software
 - What if you did? Can you undo a committed database transaction?

Error Handling - Issues in Real World

- Consider a situation where an error happens today, but it requires you to 'undo' transactions that took place over the last 2 weeks
 - This happens more often than you think when you update large software systems.
- Any application that has long lived work-flows and complex state can have this type of problem.
 - Things get even more messy since the data is often not reliable (or) unavailable.
 - Example: Update to a student pre-requisite checking system where rules change, but older rules were incorrectly recorded.

Software Deployment - Adapt




Figure 1: Activities of the Software Deployment Process.

<https://www.ics.uci.edu/~andre/papers/T3.pdf>

Adapt

- Adaptation is the ability of the software to adapt to an external change
- Potential aspects that can change:
 - External hardware and associated drivers (Video card or Sound driver)
 - External components (typically operating system update or web services that are used)
- Adaptation is generally triggered as part of the activation step
- Adaptation may trigger an update for new components (or) may need a fresh install

Software Deployment – De-Release




Figure 1: Activities of the Software Deployment Process.

<https://www.ics.uci.edu/~andre/papers/T3.pdf>

De-Release

- The process involves slowly removing support services and formally retiring a product
- Typically triggered when a new release is created, a previous release may be retired
- Most of the time this process involves advertising:
 - Removal of support services for a particular version of a product or a product line
 - Removal of development resources
 - Forcing licenses to expire
- The actual software itself is not formally removed

Timed Licensing

- Some software systems are released under a timed license
 - You lease the software for a period of time (e.g. 12 months)
 - After this period, the software will cease to function
- This model allows for a more effective de-release
- Commonly applied in large enterprise systems
 - Software running mainframes for instance.

Transactional Integrity

- The following activities must maintain transactional integrity and may need authorization:
 - Install, Update, Adapt and Activate
 - De-activate and De-install
- Ideally should be able to compensate for:
 - Power-failure
 - Component failures
- Must allow the user to return to previous state (in worst case scenario)
- Good improvements in Windows 7, Win 8, etc ...

Data Migration

- Data migration is often needed when,
 - Installing over a previous version
 - Updating components
 - Adapting to changing external configuration
- Some types of migration can be undertaken by the configuration step (either during Install or Update)
 - A database schema change can be performed by running a set of 'SQL' scripts and transferring data
- Some migration may require activation

Short Problem 5

- You have an application that has over 12,000 tables in the database. A major revision has created another 1,500 new tables and it also involves moving existing data between tables.
- What approach will you take for the data migration when you update the software? What are the risks in this approach?

Lecture Overview

- In this lecture we will cover,
 - Un-installing software
 - Updating and Adapting
 - Activating and De-activating
 - **Deployment Practices and Tools**
 - **Multi-tier and Web Applications**

Deployment Practices - I

- The following are certain deployment practices that should ideally be supported in good software:
 - Check configuration before install
 - Automatic resolution of dependency issues
 - Automatic (with consent) update process
 - Rollback of an update is possible
 - Rollback of an install is possible
 - Install and Update requires minimal/no downtime
 - Deployment configurations for test, dev, staging, live environments

Deployment Practices - 2


- Data updates are possible and do not require downtime (e.g. for anti-virus definitions)
- Logging of all transactions (install, update, adapt, activation, de-activation, de-install)
- Backups are automatically performed into an archive if product is removed (or) updated (or) upgraded software is able to monitor for a new version and inform user
- Multiple delivery models are available (Internet, DVD, phone)

Deployment Tools

- Some tools available to help automate deployment:
 - Windows Installer XML (WiX)
 - Chocolatey
 - RPM
 - NSIS (Nullsoft Scriptable Installation System)
 - New ones ...
- These tools support the deployment model to various levels -- some aspects are not covered by the tools (e.g. retire, activation and deactivation)

Deployment - Tiered Applications


- Deployment may not update all tiers at the same time
- DB server may be updated independent of application



<https://docs.oracle.com/cd/E19263-01/817-5764/architecture.html>

Deployment - Tiered Applications


- Windows installer / RPM tools do not work well in these environments
- Custom deployment scripts need to be developed



<https://docs.oracle.com/cd/E19263-01/817-5764/architecture.html>

Middleware


- Many Web applications make use of middle-ware to communicate between layers
- Middleware has to be installed and configured to work together



<https://docs.oracle.com/cd/E19263-01/817-5764/architecture.html>


Deployment for Larger Systems

- The number of components that need to be carefully considered increases
- Current tools not mature for installations spanning many tiers -- but custom scripts can be created



Rich Web Applications

- The web clients can now be Flash, Silverlight, WPF as well as HTML5/CSS/javascript -- this adds additional complexity
- The interfaces between tiers is often fire-walled
- Minimising downtime is a challenge -- esp. when many physical machines need to be updated
- Can be difficult if update needs re-booting and systems do not restart properly



Minimising Downtime - Web Applications

- Large web applications achieve this by running parallel infrastructures
- New version is deployed to a “new live environment”
 - New user logins are directed to this infrastructure
 - Data migration is triggered at login time using specific data migration APIs
- User sign-off from old infrastructure (or) eventually their sessions will expire
 - Over a couple of weeks, the old infrastructure will have no further users and can be retired

Minimising Downtime - Web Applications

- Parallel infrastructures are expensive to run -- but possible if you have the user base to cover costs
- Product design implications:
 - It should be possible to perform data migration quickly and automatically -- must be engineered into the product
 - Functions need to be transactional and stateless for best results -- works where sessions eventually expire like for Web applications
- Partial updates (except for security patches) are generally avoided to infrastructure

Web Applications

- Activation in these application involves an additional aspect
 - The standard start-up process still applies
 - The login/sign-on process is treated as an activation of the application
 - Authentication check is performed
 - Data integrity checks are also triggered to help with migration of any data after update

Software Deployment Model




Figure 1: Activities of the Software Deployment Process

<https://www.ics.uci.edu/~andre/papers/T3.pdf>

- NEXT ...

Deployment - Story so far..

- The software deployment model
- Covers activities and relationships between these activities
- Work items inside each activity and issues that need attention
- The model broadly covers “what”
- We did not cover - how to execute the activities nor large scale environments
- Next, we will look at methods to deploy software

Related Learning objectives

1. Describe the activities in software deployment and apply these in a problem context.
2. Apply knowledge of software environments to plan development and deployment.
3. Describe and classify issues that drive software maintenance.


Lecture overview

This lecture will cover,


- **A motivating scenario (Jakarta EE Application)**
- Build Numbering and Management
- Deployment Approaches
 - Manual, Scripted, Language and Model-Based

High Level Context Diagram (Example)

- How do we plan for a deployment into this environment?



Real World System (Example)



ESB - Enterprise Service Bus, SES - Security Enforcement Service

It has a mix of .NET and Jakarta EE applications,
Different DBMS as well as Security Services

Jakarta EE Deployment Scenario

- Lisa has to deploy a multi-tier JEE application:
 - MySQL Database server 5.0.x (clustered)
 - Active MQ (Message Queue)
 - JBoss Application Server
 - Spring Framework
 - 15 different java libraries for a range of tasks
 - Apache Web server
 - Monitoring and Administration applications
 - F5 load balance
 - LDAP server for authentication

Jakarta EE Deployment Scenario

- Staging/Live Environments:
 - 2 server grade machines for DB, 2 machines for web servers, and 1 machine for Application server
 - 1 desktop grade machine for monitoring and administration
 - All machines allow remote access from certain IP addresses (inside the DMZ)
 - All servers use Windows Server O/S

Jakarta EE Deployment Scenario

- Test Environment
 - Database, Web server and Application server all run on a single physical server grade machine
 - 5 testers each with desktop machines use the same test server
- Development Environment
 - Each developer has a high-end desktop (Windows 10) -- 10 developers
 - Git, Defect logging and Integration service runs on a single Linux server
- Integration service shares database

Jakarta EE Deployment Scenario

- Development schedule is 6 week time-boxed iterations
 - Team delivers a new build ready into Staging every 6 weeks
 - Team delivers and deploys a new build ready for Testing every fortnight
- Business decision causes one of the staging builds to be pushed into a live environment as an upgrade
- Deployment must support *full rollback* on all environments and must have an automated smoke test (i.e. to check everything works as expected)

Situation...

- Every year we need to undertake:
 - 26 deployments into test environment (with a rollback option)
 - 8 deployments into staging environment
 - Probably 4 deployments into live environment (certainly no more than 8)
 - Update Windows O/S 12 - 15 times (on average Microsoft releases an update once a month)
 - Update/Patch all other software in use (typically no more than once a month)
 - One major upgrade of some software systems

Lecture overview

This lecture will cover,

- A motivating scenario (Jakarta EE Application)
- **Build Numbering and Management**
- Deployment Approaches
 - Manual, Scripted, Language and Model-Based
 - We are taking a short de-tour, but will return back to the Jakarta EE scenario after this topic.

Build Management (Jakarta EE Scenario)

- Given the pace of development and the number of releases into various environment -- how does one keep track of things?
- Current good practice in these scenarios is to use Incremental single build numbering for Test/Staging and Live environments
 - Builds are numbers 1, 2, 3 ... 1341, 1342
 - Environment is pre/post fixed to build
 - Example:T3987 (Test build 3987), S3991 for staging, L3991 for Live

Builds, Versions and Branches ...

- Incremental build numbering is easy to use -- especially when logging defects
 - Works quite well for Test/Staging/Live environments
 - Developers can still work with internal branches etc. -- but any build created for testers/QA/customers is tagged with a simple incremental build number
- Traditional version numbering (e.g. 5.0) is assigned to a build by Product Managers or Sales/Marketing

Builds and Revisions

- Incremental build numbers work well in most cases -- but have limitations for 'product lines'
- Example Situation:
 - Let's say, Build 2345 of Windows 8 goes Live
 - Revisions to Win 8 will be released for testing as 2345.1, 2345.2 ... (till product line is retired)
- When a new major product is started (say Windows 10) -- they seed the base code into the repository, prepare a Build 0 and start again

Builds – Revisions and Defects

- Consider this situation:
 - Windows 8 Build 2345.581 fixes serious flaws
 - Windows 10 build numbers are based on the year and month (1903=March 2019).
 - How do we ensure that Windows 10 code also fixes these serious issues?
- Solutions:
 - Merge code from Windows 8 to Windows 10: May work, but generally not practical
 - Realistically, Windows 10 testers have to replicate and add this to defect log with note on Windows 8 fix.

Builds - Revisions - Defects - Merge

- Merging code fixes from a branch into new product line may work -- only if this does not cause more problems
- In practice,
 - Testers will replicate and log defect for new system, may add a note about occurrence in Windows 8 (adds cost as we have to log defect twice)
 - Developers are left with choice to look at the fix and use that code (or) patch it again
- This is a trade-off - often faster to put in a *kludge* (workaround).

Back to Jakarta EE Scenario ...


- In typical well managed enterprise scenarios
 - Incremental build numbering (some use ISO date format -- e.g. 20150826)
 - Minimum use of branches with revisions
- Build numbers are not just for the 'software being developed*' -- they encapsulate a bit more
 - The exact version of all components, including the operating system, database server in use
 - This level of control is only found in mature organisations (as it adds cost)

Jakarta EE Scenario

- What are we dealing with?
 - Multiple environments (Dev/Test/Staging/Live)
 - Need to have the ability to rollback changes
 - Will need to migrate data
 - Relatively high-frequency deployment to test environment
 - Very low frequency live deployment
 - Staging environment may not mirror live (esp. if build is not pushed across)
 - Stable development environment

Jakarta EE Scenario

- How do we deploy the software and all associated components?
 - Assume Operating System is installed -- but nothing much else is configured/setup
- Options are:
 1. Write steps, manually execute deployment and manually verify configuration after activation
 2. Partial automation -- using scripts (Ant and Python)
 3. Use a deployment language
 4. Model environment and use model based deployment



Next week...

- Deployment scripts and models in more detail